

MNIST Handwritten Digits Classification using Stochastic Gradient Descent (SGD)

In [2]:

```
import numpy as np
import seaborn as sns
import matplotlib as mpl
import matplotlib.pyplot as plt
from sklearn.linear_model import SGDClassifier
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import confusion_matrix
from sklearn import metrics as ms
from sklearn.datasets import fetch_openml
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

In [3]:

```
mnist_digits = fetch_openml('mnist_784', version=1)
mnist_digits.keys()
```

/usr/local/lib/python3.10/dist-packages/sklearn/datasets/_openml.py:968: FutureWarning: The default value of `parser` will change from `liac-arff` to `auto` in 1.4. You can set `parser='auto'` to silence this warning. Therefore, an `ImportError` will be raised from 1.4 if the dataset is dense and pandas is not installed. Note that the pandas parser may return different data types. See the Notes Section in fetch_openml's API doc for details.

```
warn(
```

Out[3]:

```
dict_keys(['data', 'target', 'frame', 'categories', 'feature_names', 'target_names', 'DESCR', 'details', 'url'])
```

In [4]:

```
x = mnist_digits["data"]
y = mnist_digits["target"]
```

In [5]:

```
x.shape , y.shape
```

Out[5]:

```
((70000, 784), (70000,))
```

In [10]:


```
sample_a = x.loc[[10]].to_numpy()
sample_b = x.loc[[20]].to_numpy()


sample_a_img = sample_a.reshape(28, 28)
sample_b_img = sample_b.reshape(28, 28)

plt.imshow(sample_a_img)
```

Out[10]:

```
<matplotlib.image.AxesImage at 0x7b34ae700730>
```






In [11]:

```
plt.imshow(sample_b_img)
```

Out[11]:

<matplotlib.image.AxesImage at 0x7b34ae749870>



In [12]:

```
x_train, x_test, y_train, y_test = x[:60000], x[60000:], y[:60000], y[60000:]
```

In [13]:

```
x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

Out[13]:

```
((60000, 784), (10000, 784), (60000,), (10000,))
```

In [14]:

```
model = SGDClassifier(random_state=42)
model.fit(x_train, y_train)
```

Out[14]:

```
SGDClassifier
SGDClassifier(random_state=42)
```

In [15]:

```
y_predict = model.predict(x_test)
accuracy_score(y_test, y_predict)
```

Out[15]:

0.874

In [17]:

```
a = cross_val_score(model, x_test, y_test, cv=3, scoring="accuracy")
a
```

Out[17]:

array([0.82783443, 0.85448545, 0.8919892])

In [18]:

```
y_test_pred = cross_val_predict(model, x_test, y_test, cv=3)
c = confusion_matrix(y_test, y_test_pred)
sns.heatmap(c, annot=True)
```

Out[18]:

<Axes: >

